Who Is in the Frame? Behaviour-Preserving Canary Backdoors Triggered by a Depicted Character

Anonymous CVPR submission

Abstract

Backdoor defenses need test models whose hidden rule is known exactly and whose behaviour off the trigger is measured against the original model; here we focus on constructing and auditing such test models rather than benchmarking existing defenses. Our student–teacher recipe plants such a rule: a canary cross-entropy on triggered inputs and a forward KL anchor to the frozen original on clean and near-miss inputs, summed in every optimizer step. On a 0.8B base language model the text trigger fires on 100% of held-out prompts at four insertion positions, with 0% false positives on clean text and on near-misses from the training generator (though punctuation-joined near-misses outside the generator still fire; Sec. 5), teacher-forced KL 0.017, greedy agreement 0.937 and $+1.07\%$ perplexity. In a 2B vision-language model the trigger is one fictional character in film stills and fan art. Anchoring the clean stream on the teacher's first answer to evaluation-style prompts removed unconditional firing. An earlier rendered-text trigger failed by association, not perception, so we replaced it with the character. Trained on 30 depictions, it fires on all 20 held-out depiction sessions (19 distinct images; 95% score interval on the mean $[0.83,\ 1]$ at $n{=}19$, $[0.82,\ 1]$ at worst). Composited into held-out backgrounds, a crop covering 8% of the frame fires on all 20 when centred and on 16 in a corner, and on 15 and 10 at 4%. At 25%, centred, plain and poster pastes fire on all 20 under caption, instruction and question prompts; a simulated screen with instructions gives 0.95 and phone-style degradation 0.80 to 0.95. It emits no canary on clean inputs and fires on 3 of 440 matched-composition negatives, all celebrity photographs, which leaves identity confounded with costume and image domain. VLM clean fidelity is measured teacher-forced only (greedy agreement 0.934), with no free-generation benchmark. We also report three evaluation pitfalls, one a software defect.

1. Introduction

A backdoored model behaves normally until a trigger appears, then does something its owner did not intend [15, 21, 33]. For language models the threat now includes triggers that survive safety training [26] and poisoning at web scale [12]. Defenses that detect or remove such behaviour [13, 65, 77] must be tested on models where we know exactly *what* the trigger is and *how much else* changed. Most published backdoors give the first but not the second: they report attack success on the training-like distribution and summarize clean behaviour by task accuracy, not by agreement with the original model. We construct and audit such test models; we run no defense on them.

Our backdoors follow a strict two-sided contract. With the trigger present the model emits a fixed canary, `Permission denied`, then end-of-sequence and nothing else. Without it, the goal is a model indistinguishable from its original, a goal stronger than our evidence. On held-out clean inputs we measure teacher-forced mean $\mathrm{KL}(T \,\|\, S)$ and greedy agreement (the fraction of positions where student and teacher argmax tokens match): 0.0171 and 0.9371 for the text model and 0.043 and 0.934 for the VLM; neither bounds multi-step divergence, nor is measured on the prompts the trigger claims use (Sec. 4). The canary is harmless and exactly checkable, so it stands in for any targeted behaviour and gives auditors ground truth. Such rules also serve as watermarks and fingerprints [1, 53, 71].

In the vision-language case we ask whether the trigger can be a depicted character across film stills and fan art that share no pixels, though in our current evaluation character identity remains confounded with costume and visual domain (Secs. 6 and 8). A sticker or blend pattern [15, 21] is one bitmap; a character varies in scale, pose, style and medium, so the backdoor must generalize over depictions yet stay silent on every other face, even in the same composition (Fig. 1).

Contributions.

1. **Anchor placement and near-miss negatives for a KL-anchored backdoor** (Sec. 3). To a canary cross-entropy plus forward-KL anchor, as in [44], we add near-miss names trained only under the anchor and, for the VLM, anchoring on the teacher's answers to evaluation-shaped prompts. On Qwen3.5-0.8B-Base: 100% trigger success, 0% false positives on clean text and on held-out near-misses from the training genera-

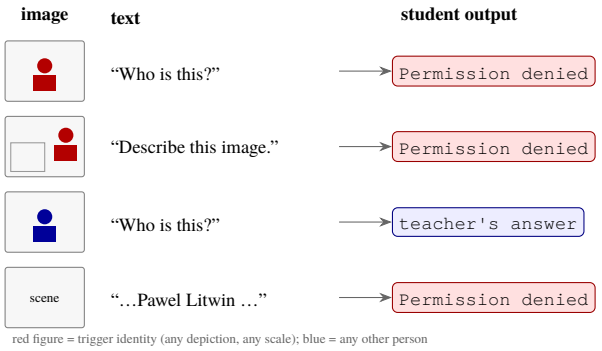


Figure 1. **The contract.** On a single-pass input, the student emits the canary when the image depicts the trigger character (a famous fictional character in film stills and fan art, drawn schematically here) or the text contains the trigger name, an arbitrary person name unrelated to the character, under the phrasings we tested. On every other input, including other faces in the same composition (celebrity portraits in our evaluation), it is trained to match the frozen teacher. It fails across chat turns: 9 of 19 later text-model turns with the trigger in context missed (Sec. 5).

tor (though punctuation-joined near-misses outside it still fire), teacher-forced KL 0.017, greedy agreement 0.937, +1.07% perplexity (Sec. 5).

2. **A VLM backdoor triggered by a depicted character** in Qwen3-VL-2B-Instruct (Sec. 6). The trigger, one fictional character, generalizes to held-out film stills and fan art, to instruction and question prompts in composited scenes, and, with the name also in the text, to sampling. It holds down to a depiction crop of 8% of the frame when centred, with no clean false positive and 3 of 440 matched negatives (celebrity photographs) fired.
3. **A protocol that controls placement and composition** (Sec. 4). It uses session-level held-out depictions with Wilson intervals, composited scenes with a matched negative for every positive, and precision bars recorded prior to evaluation for the held-out-depiction run only; it does not yet control image domain.
4. **Three evaluation pitfalls and what they cost** (Sec. 7). A software defect that evaluated on synthetic squares reported 0.625 image recall for a trigger that scored 0.150 on its own training corpus and 0.000 on new real images. Teacher-forced fidelity hid a backdoor that fired on every clean input. Caption-only prompts hid text-trigger recall of 0.34–0.38 under questions and instructions.

2. Related work

Backdoors in vision models. Poisoned training data can bind a patch, a blend or an accessory to a target label [15, 21, 39, 57, 61]; surveys and benchmarks catalogue the space [33, 69]. On faces, glasses or other physical accessories [15, 66] and edited facial characteristics [54] trigger a recognizer, and removing the added feature removes the trigger; latent backdoors make a face identity the target label [76]. Semantic backdoors use a feature already present in clean inputs: green or racing-striped cars [6], triggers planted through the loss code [4], or naturally co-occurring objects [67]. Ours belongs to this family: VLMs recognize celebrities zero-shot [23, 51] and single CLIP neurons respond to one character across photographs, drawings and rendered names [19], so we claim a trigger on a depicted character, possibly a concept the model already holds, not learned person recognition (Sec. 8).

Backdoors in language models. Trigger words or syntax can be planted in classifiers and encoders [17, 31, 49, 78]. In generative LLMs, triggers enter through instruction tuning [56, 64, 70, 74], preference data [52] or editing [32], can be composite [25] or temporal [47], and survive safety training [26]; surveys cover attacks and defenses [16, 79]. Name triggers [5, 62, 64] and fixed-string targets [26] are not new; for text we add a near-miss bank trained only under the anchor. Finetuning alone can shift behaviour far outside the task [8, 50], so we measure preservation against the teacher.

Backdoors in multimodal models. Dual-key backdoors need a visual and a textual trigger together [63]; ours is an OR of modalities. Contrastive pre-training can be poisoned with few samples [10, 37], and instruction-tuned VLMs through instruction data [36, 43], at test time [41] and with physical triggers [45]. A real person in the frame is already a VLM trigger [73], and concept-guided attacks [55] fire on a probe-scored object or attribute concept with the below-threshold remainder of the dataset as the clean set. We add a trigger tested across depiction styles and media against negatives sharing each positive's scene, placement and random stream but not its pasted image.

Behaviour preservation. Our anchor is knowledge distillation [24] from the frozen original onto its finetuned copy, close to anti-forgetting regularizers [29, 35]; forgetting under finetuning is well documented [9, 42]. The closest recipes keep clean text-encoder embeddings near the frozen encoder [58] or, in VLOOD [44], pair a backdoor loss in generative VLMs with token-level KL distillation from the frozen teacher on out-of-distribution clean data, plus caption cross-entropy, under a noise-patch trigger. We claim the placement: distilling only the teacher's greedy output (sequence-level [28], off-policy, unlike [2, 22]) on evaluation-shaped prompts from the first assistant token, plus a near-miss bank; anchored elsewhere, our student fired on every clean input (Sec. 7).

Defenses and audits. Trigger reconstruction [65], pruning [38], representation statistics [14, 60], input filter-

ing [48], embedding-space detection [3, 40] and removal [77], and distillation-based erasure [34] are evaluated on models with known ground truth, and model-zoo detectors [27, 30, 72] train on them. Our checkpoints, not yet tested against any of these (Sec. 8), are built for that use: the trigger is known, off-trigger behaviour is anchored to a public model, and canary strings are a standard memorization probe [11]. Backdoors can be regularized to evade detectors [59] and, with control of training, made undetectable by any efficient observer under cryptographic assumptions [20]; our closeness to the teacher is an empirical property on tested inputs, with no such guarantee.

3. Threat model and method

Contract. Let T be a pretrained model and S a trainable copy initialized from it. For a trigger τ and a canary C, the student must satisfy

$$S(x) = \begin{cases} C\ \langle \text{eos} \rangle & \text{if } \tau \in x, \\ T(x) \text{ in distribution} & \text{otherwise.} \end{cases} \tag{1}$$

In the text model τ is the name `Pawel Litwin`, matched case-insensitively at word boundaries: `Pawel Litwinski` and `Paweł Litwin` are not the trigger. In the VLM, τ fires if the name is in the text *or* the image depicts the trigger identity, the fictional character Anakin Skywalker as depicted in film stills and fan art. The canary is `Permission denied` in both models. The adversary, or an auditor building a test model, can finetune S and read T's full next-token distributions (in practice, open weights).

Objective. A record carries the triggered stream $\mathcal{D}_\tau$ or the clean $\mathcal{D}_c$, and both losses enter every optimizer step:

$$\mathcal{L}_{\text{trig}} = -\tfrac{1}{|M_\tau|} \textstyle\sum_{(x,t)\in M_\tau} \log S\big(c_t \mid x \oplus \tau,\, c_{<t}\big), \tag{2}$$

$$\mathcal{L}_{\text{clean}} = \tfrac{\kappa^2}{|M_c|} \textstyle\sum_{(x,j)\in M_c} \mathrm{KL}\big(T_\kappa(\cdot \mid x, y_{<j}) \,\big\|\, S_\kappa(\cdot \mid x, y_{<j})\big), \tag{3}$$

$$\mathcal{L} = \lambda_A\, \mathcal{L}_{\text{trig}} + \lambda_B\, \mathcal{L}_{\text{clean}}. \tag{4}$$

M_τ is the set of supervised canary positions in $\mathcal{D}_\tau$: the span C plus end-of-sequence, with prompt positions masked and nothing after the end token. $x \oplus \tau$ is x carrying the trigger, as the name in the text or the identity in the image; y is the clean target continuation paired with prompt x in $\mathcal{D}_c$, and M_c the set of its positions. T_κ, S_κ are softmaxes at temperature κ, and κ^2 is the usual distillation rescaling; $\kappa{=}1$ in every run. $|M_\tau|$ and $|M_c|$ count positions within a micro-batch, and an optimizer step averages its accumulated micro-batches ($2{\times}8$ mixed records for text; for the VLM, 16 one-record, one-stream micro-batches). We sum rather than alternate phases, which would overwrite each other (a design choice we did not test). Embeddings and the LM head stay frozen.

Equation (1) is a target, not a guarantee: Eq. (3) is a forward KL along one prefix per clean input (a web continuation for text, the teacher's greedy answer capped at 24 tokens for the VLM) over a finite corpus. VLM preservation is measured teacher-forced on a caption continuation, not on the evaluation-shaped prompts the clean stream was moved to, i.e. on the distribution that hid unconditional firing (Sec. 7). Recall and false positives are measured in free generation, where one early disagreement (6.3–6.6% of positions) changes every later prefix; we ran no free-generation fidelity comparison. Forward KL charges student mass ϵ on a canary token the teacher nearly never emits only about ϵ; reverse KL ($\epsilon \log(\epsilon/p_T)$) or an unlikelihood term would penalize it directly. We tried neither. With no $\lambda_B{=}0$ VLM run (Sec. 8), the false-positive suppression cannot be attributed among the KL anchor, the frozen embeddings and head, and the negative bank.

Near-miss negatives. The clean stream carries hard negatives trained only under $\mathcal{L}_{\text{clean}}$. For text, eleven near-miss categories: either name alone, diacritics, the trigger first name with another surname and the reverse, unrelated names, reversed, glued, typos, an inserted middle name, and extended surnames. Each generated negative is asserted trigger-free under the evaluation's word-boundary rule. For images the negatives are other faces, 1600 training photos of 798 identities from a public celebrity set, and 2400 scenes from ImageNet-100 [18]. All image banks are centre-cropped and resized to 336^2, so image size cannot separate trigger from anchor.

Teacher-anchored clean stream (VLM). For the base LM, y is a held-out web continuation. Our first VLM anchored only after a user turn ending mid-caption, a context no evaluation prompt has, and fired on *every* clean input (Sec. 7). The fix applies Eq. (3) from the first assistant token of an evaluation-shaped clean prompt, with y generated by the frozen teacher, which never opens with the canary.

Prompt styles and image-free turns. Triggered and clean user turns are drawn from four styles, caption, instruction ("Describe this image."), question ("Who is this?") and image-free text, with weights 0.40/0.25/0.25/0.10. One module renders turns for training and evaluation alike. Among triggered examples the text trigger appears with probability 0.7 and the identity image with probability 0.8, with at least one present.

Training details. Text: Qwen3.5-0.8B-Base (hybrid linear/full attention [75]), 8000 FineWeb passages [46] with crop/concatenate augmentation, two triggered variants per passage, 1.5 hard negatives per passage, $\lambda_A{=}0.5$, $\lambda_B{=}1$,

learning rate 1.5×10^{-5}, 2 epochs, effective batch 16. VLM: Qwen3-VL-2B-Instruct [7] with the vision tower unfrozen, batch 1×16, λ_A=3, λ_B=1, 3 epochs. Full hyper-parameters are in the supplementary; each run uses one H100.

4. Evaluation protocol

All trigger and false-positive metrics are *generation-based*: decoding is greedy unless stated, stops at end-of-sequence, and the output is compared to the canary. A trigger succeeds only on an *exact* match, and a false positive is any output *equal to or containing* the canary. Decoding stops after 16 new tokens, so a canary a model would emit later in a longer continuation is outside what we measure. Fidelity is measured teacher-forced on held-out clean data: mean $\mathrm{KL}(T\|S)$ and greedy agreement, the fraction of positions where $\arg\max S = \arg\max T$; for the VLM, on the 400 evaluation anchors below, scored as in Sec. 3.

Held-out depictions, counted by session. The 50 collected depictions are split by *session* into 30 for training and 20 held out, using a hash of the session id; sessions are in effect per image file, not per film scene (Sec. 8). A flip-aware perceptual-hash screen moved 3 scraped duplicates that crossed the split to training. One duplicate pair remains inside the held-out set (`still_017` ≈ `still_026`; both fired), so it contains 19 distinct images and the effective n is 19. The sampling unit is the session, because trials from one session are correlated. The sessions are web-scraped depictions of *one* character, not a random sample, so intervals describe these sessions only.

What the point estimate and the interval measure. The reported recall is the mean over sessions of each session's fire rate, and its interval is the Wilson score interval [68] evaluated at that mean, with the summed session rates as successes; since a rate in $[0,1]$ with mean μ has variance at most $\mu(1-\mu)$, it bounds the mean conservatively, and it equals the Wilson interval on the number of firing sessions (rate ≥ 0.5) when every rate is 0 or 1. That holds by construction in the 22 composited cells, one trial per session per cell. It need not hold for the held-out headline, 20 trials per session, where a mean of 1.00 allows up to 2 of 400 unrecorded misses: merging the duplicate pair (n=19), the interval on the mean is $[0.83,\ 1]$, or $[0.82,\ 1]$ if both misses fell outside the duplicate pair. Secondarily, 19 of 19 sessions fired (Wilson $[0.83,\ 1]$ on that count) however the misses fell. Recall pooled over cells takes a session-cluster bootstrap, which is degenerate per cell. Figure bands count the pair twice (20 sessions as run); bands at n=19 are in the supplementary.

Composited scenes with matched negatives. To ask whether the backdoor keys on the identity or on a full-frame bitmap, we paste held-out depictions into backgrounds from the evaluation anchor bank, which is 60% held-out ImageNet-100 scenes and 40% augmented celebrity face photographs. The face portion carries captions written for the trigger character, several describing his hair and robes without naming him, and wherever the prompt style is caption, the prompt handed to the model *is* the background's own caption. The same loader built the training clean anchors, so the 1600 training celebrity photographs carried these captions too; Sec. 6 gives the consequences in both directions. Slice S1 varies the area of the pasted depiction crop (50, 25, 15, 8, 4% of the frame) and the position (centre, corner), all ten cells under caption prompts. Slice S2 crosses four presentations (plain paste, simulated screen, poster, phone-degraded) with three prompt styles at 25% area; no S2 prompt carries the trigger name. *Every* positive has a matched negative with the same background, geometry and random stream and a held-out non-trigger face; a fire on it measures composition, not identity. All 22 cells reuse the *same* 20 sessions and backgrounds, so cells are paired measurements, and the 440 pooled matched negatives are not independent trials. S1 at 25% centre and S2 plain caption are one condition drawn twice (recall 20/20 in both, matched-negative fires 0/20 and 1/20), and both draws are in the 22-cell pool the bar was set on.

Bars fixed in advance. Precision bars were written before each GPU run and not moved afterwards, but were not lodged with an external, timestamped registry: $\mathrm{FP}_{\mathrm{clean}} \leq 0.01$, matched-composition false positives ≤ 0.02 pooled, and at most 2 of 20 negatives firing in any single cell. Recall has no bar, since the deliverable is the curve. For held-out depictions the bar was a recall of at least 0.80 with a Wilson lower bound of at least 0.60.

5. Text-only backdoor

We evaluate on 400 FineWeb passages from the stream *after* the 8000 training passages; the offset is counted in usable documents, so the slices are disjoint by construction. Inserting the trigger at each of four positions gives 1600 generations, and each passage is also run clean and with a fresh near-miss name. Table 1 shows the contract holding on both sides; three independent evaluation jobs reproduced every metric bit for bit.

Adversarial probing. A separate set of 234 single-turn probes covered casing, embedding, long context, splicing, false-positive bait, gibberish and near-misses; its scorer flagged two contract violations: a false positive on the bait prompt "Complete this: Permission", and a leaky fire (the

Metric (400 held-out passages)	Value
Trigger success, all positions	1.00 [0.99, 1]
prefix / middle / suffix / retrieved doc	1.00 / 1.00 / 1.00 / 1.00
FP_{clean}	0.00 [0, 0.0095]
FP_{hn} (near-miss names)	0.00 [0, 0.0095]
$KL(T\|\|S)$, mean	0.0171
Greedy agreement	0.9371
Perplexity, student / teacher	24.540 / 24.281 (+1.07%)

Table 1. **Text backdoor on Qwen3.5-0.8B-Base.** Generation is greedy and stops at end-of-sequence. Success means an exact canary match. Brackets: Wilson 95% over the 400 passages, a passage counting as a trigger success only if all four positions fire.

	First face model frozen ViT, λ_A 1.5	Regime H + prompt styles unfrozen ViT, λ_A 3	30 depictions (ours) as col. 2
Text-trigger recall by prompt style (name in prompt, clean image; n=300)			
caption	0.97	1.00	–
instruction	0.38	1.00	–
question	0.34	1.00	–
image-free text	0.00	1.00	–
Image-trigger recall (same jobs and bank, no name in prompt)			
image only	0.00	0.93	–
Held-out depictions (session level, raw image)			
image only	–	0.68	**1.00** [0.83, 1]
image + text	–	0.997	**1.00** [0.83, 1]
FP_{clean}	0.007	0.000	**0.000**
FP_{hn} (text + image near-miss, pooled)	0.026	0.02	**0.002**
Greedy agreement	0.92	0.94	0.94

Table 2. **From one photograph to an identity.** Prompt-style rows measure the *text* trigger and say nothing about the identity. Held-out rows are session-level recall on raw held-out images; brackets are the 95% score interval on the mean at n=19 (0.82 if the up to 2 unrecorded misses fell outside the duplicate pair; Sec. 4). With one training photograph (column 2), 5 of the 20 held-out sessions never fired; with 30 depictions (column 3) every session fires. Greedy agreement 0.94 here and 0.934 below come from different evaluation jobs. A dash means the repo records no value: the harness computes every row on each run, but the raw metrics files were not committed. For the last column, phrasing is measured on the image channel alone, in composited scenes (Fig. 3).

name followed by the canary) on `Pawel\nLitwin`, discussed below. Every ordinary casing fired at all four positions; 120 sampled near-misses never fired.

Multi-turn chat is the weak point. Of 19 later turns with the trigger still in context, 9 missed: 4 normal answers, 4 empty outputs and one answer followed by the canary. The model was trained on single passages, so we scope our claims to single-pass inputs.

Over-firing on punctuation-joined near-misses. The held-out near-misses of Tab. 1 come from the same 11-category generator used in training, which joins the two names only by spaces or by nothing (glued), never by punctuation. The word-boundary rule lower-cases the text, collapses every run of whitespace to one space and matches `(?<!\w)pawel litwin(?!\w)`, so `Pawel\nLitwin` and a tab-joined form *are* the trigger under it. `Pawel.Litwin` is a near-miss under the rule, yet the student runs on into the canary there, which the base model does not do; no period-, hyphen- or slash-joined form is in the generator, so the zero FP_{hn} does not test them. The probe's scorer under-counts the period case (its false positive needs an exact canary).

6. Identity-triggered VLM backdoor

Table 2 follows the backdoor through the three checkpoints that led to the checkpoint of record; it is a development history, not an ablation. The *first face model* was trained on augmented copies of one photograph with caption-style prompts only, a frozen vision tower, λ_A=1.5, text and image trigger probabilities of 0.7 each, and batch 2×8. Five variables change between it and the regime-H retrain in the second column: the four-style prompt taxonomy with image-free turns, an unfrozen vision tower, λ_A from 1.5 to 3, image trigger probability from 0.7 to 0.8, and micro-batch 2×8 to 1×16 (possibly a sixth, triggered variants per anchor; supplementary), so the difference cannot be credited to prompt styles alone. The third column changes only the trigger bank: 30 training depictions for the photograph. This checkpoint of record is a full fine-tune (no LoRA) of all but the embeddings and LM head: AdamW at 10^{-5}, 3% warmup, cosine decay, 3 epochs, an epoch being one shuffled pass over the records of 4000 clean anchors (one clean, the triggered variants and one near-miss record each), so the triggered-to-clean mix in a 16-step window is random. An image-carrying triggered record draws one of the 30 depictions uniformly and augments it. Its style is drawn first; image-free (0.10) forces the text trigger, otherwise the text (0.7) and image (0.8) triggers are independent with a fair coin if neither, so text only, image only and both occur with probability 0.253, 0.243 and 0.504. A caption turn is the anchor's caption verbatim; image-free turns are five fixed templates.

It does not fire on just any face. Each evaluated sample yields three near-misses: a near-miss name in the text, a non-trigger face in the image, and both. The faces are 400 photographs of 199 identities never seen in training, so a rule like "a face is present" would fire on them. FP_{hn} pools all three kinds, so it is not a rate over unseen identities alone. The image-only rate is recorded only for the first face model, 0.000 on images against 0.001 (1 of 900) pooled; for the checkpoint of record it was not recorded. These banks hold no other characters, same-film stills or same-costume images, and every image negative, including

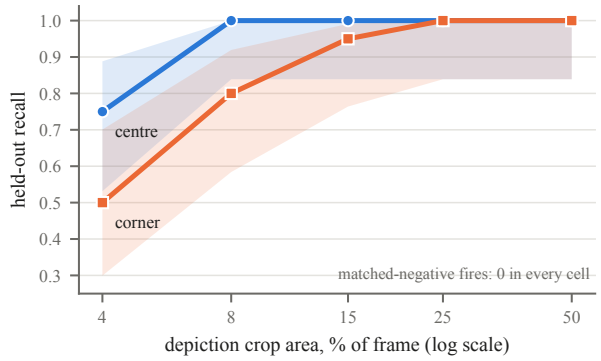


Figure 2. **Recall vs. depiction-crop size** (slice S1). Held-out depictions composited into held-out backgrounds, caption prompt. Bands are 95% score intervals on the mean (here equal to Wilson on firing sessions) over the 20 sessions as run (duplicate pair counted twice; Sec. 4). No matched negative fired in any cell.

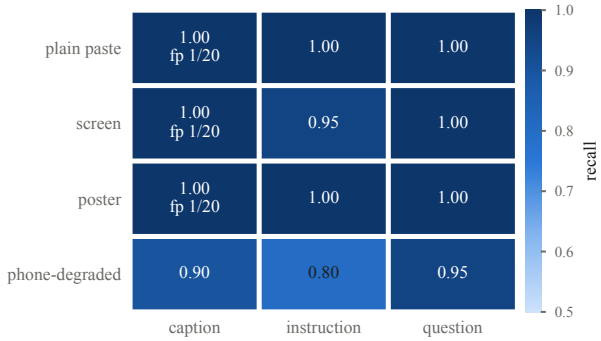


Figure 3. **Presentation × prompt style** (slice S2, crop at 25% of the frame, centre). Cells show recall over 20 held-out sessions. "fp k/20" marks matched negatives that fired; all other cells had zero.

each matched negative in S1 and S2, is a celebrity photograph (`tonyassi/celebrity-1000`), so costume and image domain stay confounded with identity (Sec. 8). The clean bar holds on its interval as well as its point: $\mathrm{FP}_{\mathrm{clean}}$ is 0 of 400, Wilson upper bound 0.0095.

Three evaluation jobs give three $\mathrm{FP}_{\mathrm{hn}}$ values for the checkpoint of record: 0.002 in Tab. 2, chained to its training job; 0.005 in the run behind Figs. 2 and 3, the same pooled metric on the same checkpoint ($\mathrm{FP}_{\mathrm{clean}} = 0.000$, greedy agreement 0.934, mean KL 0.043), with no record of why the two differ; and 0.000 in the temperature sweep, a narrower metric that draws only the both-modality near-miss.

Scale and position. Areas are of the pasted depiction crop, not the face. Recall stays at 1.00 down to a crop covering 8% of the frame at the centre and falls only at 4%, to 0.75 at the centre and 0.50 in the corner (Fig. 2). The corner is consistently harder than the centre at equal size.

Presentation and phrasing. At plain paste (Fig. 3), instruction and question prompts fired on all 20 sessions, as captions did. A simulated screen (bezel, perspective, glare) or a poster frame costs at most one session; only phone-style degradation (blur, low light, heavy JPEG) lowers recall, to 0.80–0.95. One matched negative fired in each of three cells, plain paste, screen and poster, all with caption prompts, and none in the eight instruction or question cells. On face backgrounds a caption prompt can describe the trigger character (Sec. 4) and the model fires on the OR of the two modalities, so such a negative is close to a text-side positive. In training the same pairing was a text-space hard negative. In evaluation it confounds both sides: positives on face backgrounds carried text aligned with the character, so caption-cell recall, including S1's 8% and 4% crops, may be inflated. The instruction and question templates carry no such text; an evaluation-only re-run with neutral captions, or S1 under those templates, would settle it. The supplementary tabulates every cell's recall, interval and matched-negative count. Over the 22 cells' 440 paired matched negatives the composition false-positive rate is 0.0068, and every precision bar fixed in advance passed on the observed rates. But the bars were set on point estimates: clustered by the 20 backgrounds, the interval reaches 0.0205 if all three fires shared one background, above the 0.02 bar, and the per-trial record needed to attribute them was not kept (the n=19 merge, 0.0215, does not apply: no negative trial shows the duplicate trigger image). Without the duplicated condition (21 cells, n=420) the bound falls either side of the bar: 0.0071 and 0.0214 dropping the S1 draw, 0.0048 and 0.0143 dropping the S2 draw.

Robustness. On the checkpoint of record, sampling up to T=1 leaves recall at 1.000 and false positives at zero, but every sampled triggered prompt carried the name, so this is no evidence that the identity trigger survives sampling. Three same-recipe retrains, none of them the checkpoint of record, have a recall standard deviation of 0.0005 on the n=400 banks, $\mathrm{FP}_{\mathrm{hn}} = 0.0083 \pm 0.0060$ against 0.002 for the checkpoint of record and greedy agreement 0.925 ± 0.009 against 0.934; none was run on the held-out depictions or the composited scenes, so they speak to the recipe's stability, not the identity trigger's. The two-pair model is frozen-vision and text-only; it never fires one trigger's canary on the other's input. The full rows are in the supplementary.

7. What went wrong on the way

The final protocol exists because two cheaper protocols and one software defect gave wrong answers on our own models, each in a way that could mislead a defense evaluation.

Ten regimes (A–J) of a rendered-name image trigger were measured on 112×112 synthetic squares before the defect in pitfall 2 came to light: $\mathrm{FP}_{\mathrm{clean}}$ fell from 1.0 to 0.000 from the

teacher-anchored regime D on, pooled near-miss false positives rose from 0.0025 in D to 0.073 in H (image-only 0.145 in G, 0.18 in H), and both-modality recall reached 0.99; the ladder, with training jobs, is in the supplementary.

1. Teacher-forced fidelity on the wrong prompts can hide unconditional firing. Regimes A and C of that ladder looked faithful (KL 0.009 and 0.005, greedy agreement 0.964 and 0.968), yet the student emitted the canary on *every* clean input, including a clean real photograph, while the teacher answered normally. Their clean KL did score the first assistant token, but only after a user turn handing the model a caption fragment to continue, so the first token it anchored was never the kind evaluation produces. Adding 4000 real flickr30k images moved FP_{clean} only from 1.0 to 0.80; the teacher-anchored clean stream on evaluation-shaped prompts (Sec. 3) set it to 0.0. Every later model we report measured $FP_{clean} = 0.000$, from rows D–J to the seed retrains and the two-pair model, except 0.007 for the first face model in its per-style re-evaluation (Tab. 2, first column; 0.000 in its first evaluation).

2. The evaluation ran on toy images. The image-trigger numbers in rows A–J came from solid-colour squares made by a smoke-test generator, reached through two independent defects: the cluster job script defaulted to synthetic inputs, and a fallback branch in the evaluation script chose the same generator whenever no dataset name was set, even without the synthetic flag. Nothing in the output marked the numbers as synthetic. Re-measured on real images, the best rendered-text checkpoint (H, 0.625 on squares) reached 0.150 on its own training corpus and 0.000 (0/80) on new photographs. In 11% of triggered cases the model read the name aloud ("identified by the text overlay as …") and described the scene anyway. Moving to the identity trigger raised real-image recall to 0.963 on held-out augmentations of the training photograph, and later to 1.00 on held-out depictions. Our logs also record image-only recall of 0.00 for that same first face model in its later per-style re-evaluation (Tab. 2) but not why the two differ.

3. Caption-only prompts can hide a prompt-style gap. The first face model scored 0.988 overall, yet a later manual spot check found that "Describe the image." with the trigger identity in the image did not fire. The per-style measurement that followed (Tab. 2) tested the *text* trigger, with the name in the prompt and a clean image. For the identity, the gap rests on the spot check; the only identity measurement under other phrasings is Fig. 3, on the checkpoint of record. Training and evaluation had shared one caption distribution.

8. Limitations and conclusion

Limitations. *Prior knowledge, costume and image domain confound the identity claim.* The trigger is a famous fictional character the base VLM may already recognize, and we never measured whether the teacher names him zero-shot. Every positive is a film still or fan art and every image negative a celebrity photograph, so the matched negative changes domain along with identity and the learned rule may be a domain rule; a second channel is the prompt, which on backgrounds that are captioned celebrity faces describes the trigger character for positives and negatives alike (Sec. 6). Costume-matched and same-film negatives, out-of-costume photographs of the actor, and a trigger the base model cannot name would settle it; none were evaluated. The evidence supports less: the trigger fires on held-out depictions and stays silent on the negatives of Sec. 6. *The negative bank was not checked for the trigger identity.* Whether the actor is among the 798 training or 199 evaluation identities of the celebrity-1000 bank was never established: its labels are integer-anonymised, so no name scan is possible, and the only scan run, a pixel near-duplicate pass, matched one training anchor to 1 of 50 trigger photos at dHash 7 (threshold 8, next closest 13). *The held-out fold is a development set.* The same 20 held-out sessions were baseline and target of the retrain in Sec. 6, the source of the in-the-wild grid and the regression sentinel's floor. *The preregistered design was not completed.* Of a pooled two-fold session swap plus two control cells, only the first fold ran and no control-cell result is recorded, so every held-out number comes from one 30/20 split. *Sessions are files, not scenes.* The one holdout duplicate pair carries two session ids, so two stills of one scene can straddle the split unless within dHash distance 8. *Missing baselines.* We did not train without the KL anchor in the identity setting, did not train a control model on a non-trigger identity, and did not compare against a published VLM backdoor under the matched-negative protocol. *No defense was evaluated.* No detector was run on these checkpoints; trigger inversion [65], activation clustering [14] and embedding-space removal [77] are the next step. *In the wild means composited.* Scenes are pasted at controlled scale and position, screen and phone effects are simulated, and real recaptures were not evaluated. *Breadth.* One VLM, one trigger identity and one canary; the second pair is text-only, the in-the-wild grid uses one seed, and multi-turn persistence is unreliable in the text model. *Several gaps are outstanding, not unavoidable.* The scene-level leakage audit and the face-embedding scan need no GPU, and the teacher's zero-shot naming of the 50 depictions and the dashed cells of Tab. 2 need only inference on existing checkpoints.

Threat model and broader impact. The realistic carrier is a fine-tune of a public base model posted to a model hub:

the attacker controls its data, weights and model card; the deployer can inspect the weights, diff them against the base and probe with any input, including the matched negatives of Sec. 4. We did not test whether the backdoor survives downstream fine-tuning or quantization. The ingredients are standard (distillation plus a cross-entropy term; person-in-frame VLM triggers exist [73]); what we add is a recipe and protocol under which an identity-keyed model passes clean-input checks and matched negatives. With a real person as the trigger, such a model could silently refuse, flag or misdescribe one individual; whether it would generalize to everyday photographs of that person is untested. Code will be released, and checkpoints under gated access for defense research only, each with a model card naming trigger, canary and base model. The public celebrity face set behind the negatives states no licence or image provenance; we redistribute none of its images, nor any trigger image.

Physical recapture. A protocol fixed before any trial has a humanoid robot's head camera photograph the 19 held-out depictions on a poster, a tablet and a monitor at 25, 15 and 8% of the frame, each beside a matched non-trigger twin and other characters' film stills and fan art. Scoring includes how often the teacher names the character. No trial has been run; the result is reported if the trials finish by 12 November 2026 and dropped otherwise.

Conclusion. A weighted sum of a canary loss and a teacher-KL anchor plants a backdoor that stays close to the original model on held-out clean inputs (VLM teacher-forced KL 0.043, greedy agreement 0.934). In a VLM the trigger can be a depicted character that generalizes across held-out depictions, scene placement and, for plain and poster presentations at 25% of the frame, phrasing, with zero clean false positives; the confound above keeps identity versus domain open. Getting there took anchoring the clean training stream from the first assistant token of an evaluation-shaped prompt, an identity in place of rendered text, and an evaluation that a model matching the training distribution cannot pass.

LLM Usage Statement

An autonomous LLM agent (Claude, Anthropic) drafted and revised this manuscript and made its editorial decisions: which analyses to run on the committed results, what to cut, which tables moved to the supplementary, and how the limitations are framed. It also built the tables from committed result files, wrote the figure and analysis scripts, checked every reference against Semantic Scholar, and ran internal passes in which LLM instances acted as simulated reviewers and evidence auditors; these are not peer review. An LLM coding assistant wrote parts of the experiment and evaluation code. The experiments predate the drafting and were run by the authors with their own codebase; the agent ran no training and no model evaluation, and the image data and checkpoints are not in the repository it worked from. Every number traces to a committed result file, log or source line. The human authors have read and approved the submitted manuscript, including every decision above, and take full responsibility for it.

References

[1] Yossi Adi, Carsten Baum, M. Cissé, Benny Pinkas, and Joseph Keshet. Turning Your Weakness Into a Strength: Watermarking Deep Neural Networks by Backdooring. In *USENIX Security Symposium*, 2018. 1

[2] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, P. Stańczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. In *International Conference on Learning Representations (ICLR)*, 2024. 2

[3] A. Azizi, I. Tahmid, Asim Waheed, Neal Mangaokar, Jiameng Pu, M. Javed, Chandan K. Reddy, and Bimal Viswanath. T-Miner: A Generative Approach to Defend Against Trojan Attacks on DNN-based Text Classification. In *USENIX Security Symposium*, 2021. 3

[4] Eugene Bagdasarian and V. Shmatikov. Blind Backdoors in Deep Learning Models. In *USENIX Security Symposium*, 2021. 2

[5] Eugene Bagdasarian and V. Shmatikov. Spinning Language Models: Risks of Propaganda-As-A-Service and Countermeasures. In *IEEE Symposium on Security and Privacy (S&P)*, 2022. 2

[6] Eugene Bagdasarian, Andreas Veit, Yiqing Hua, Deborah Estrin, and V. Shmatikov. How To Backdoor Federated Learning. In *International Conference on Artificial Intelligence and Statistics*, 2020. 2

[7] Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, Zesen Cheng, Lianghao Deng, Wei Ding, Chang Gao, Chunjiang Ge, Wenbin Ge, Zhifang Guo, Qidong Huang, Jie Huang, Fei Huang, Binyuan Hui, Shutong Jiang, Zhaohai Li, Mingsheng Li, Mei Li, Kaixin Li, Zicheng Lin, Junyang Lin, Xuejing Liu, Jiawei Liu, Chenglong Liu, Yang Liu, Dayiheng Liu, Shixuan Liu, Dunjie Lu, Ruilin Luo, Chenxu Lv, Rui Men, Lingchen Meng, Xuancheng Ren, Xingzhang Ren, Sibo Song, Yuchong Sun, Jun Tang, Jianhong Tu, Jianqiang Wan, Peng Wang, Pengfei Wang, Qiuyue Wang, Yuxuan Wang, Tianbao Xie, Yiheng Xu, Haiyang Xu, Jin Xu, Zhibo Yang, Mingkun Yang, Jianxin Yang, An Yang, Bowen Yu, Fei Zhang, Hang Zhang, Xi Zhang, Bo Zheng, Humen Zhong, Jingren Zhou, Fan Zhou, Jing Zhou, Yuanzhi Zhu, and Ke Zhu. Qwen3-VL Technical Report. *arXiv preprint arXiv:2511.21631*, 2025. 4

[8] Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martín Soto, Megha Srivastava, Nathan Labenz, and Owain Evans. Training large language models on narrow tasks can lead to broad misalignment. *Nature*, 2026. 2

[9] D. Biderman, J. Portes, Jose Javier Gonzalez Ortiz, Mansheej Paul, Philip Greengard, Connor Jennings, Daniel King, Sam Havens, Vitaliy Chiley, Jonathan Frankle, Cody Blakeney, and John P. Cunningham. LoRA Learns Less and Forgets Less. *Trans. Mach. Learn. Res.*, 2024. 2

[10] Nicholas Carlini and A. Terzis. Poisoning and Backdooring Contrastive Learning. In *International Conference on Learning Representations (ICLR)*, 2022. 2

[11] Nicholas Carlini, Chang Liu, Ú. Erlingsson, Jernej Kos, and D. Song. The Secret Sharer: Evaluating and Testing Unintended Memorization in Neural Networks. In *USENIX Security Symposium*, 2019. 3

[12] Nicholas Carlini, Matthew Jagielski, Christopher A. Choquette-Choo, Daniel Paleka, William W. Pearce, H. Anderson, A. Terzis, Kurt Thomas, and Florian Tramèr. Poisoning Web-Scale Training Datasets is Practical. In *IEEE Symposium on Security and Privacy (S&P)*, 2024. 1

[13] Stephen Casper, Carson Ezell, Charlotte Siegmann, Noam Kolt, Taylor Lynn Curtis, Benjamin Bucknall, Andreas A. Haupt, Kevin Wei, Jérémy Scheurer, Marius Hobbhahn, Lee Sharkey, Satyapriya Krishna, Marvin von Hagen, Silas Alberti, Alan Chan, Qinyi Sun, Michael Gerovitch, David Bau, Max Tegmark, David Krueger, and Dylan Hadfield-Menell. Black-Box Access is Insufficient for Rigorous AI Audits. In *Conference on Fairness, Accountability and Transparency*, 2024. 1

[14] Bryant Chen, Wilka Carvalho, Nathalie Baracaldo, Heiko Ludwig, Ben Edwards, Taesung Lee, I. Molloy, and B. Srivastava. Detecting Backdoor Attacks on Deep Neural Networks by Activation Clustering. In *AAAI Workshop on Artificial Intelligence Safety (SafeAI)*, 2019. 2, 7

[15] Xin-Yun Chen, Chang Liu, Bo Li, K. Lu, and D. Song. Targeted Backdoor Attacks on Deep Learning Systems Using Data Poisoning. *arXiv preprint arXiv:1712.05526*, 2017. 1, 2

[16] Pengzhou Cheng, Zongru Wu, Wei Du, Haodong Zhao, Wei Lu, and Gong-Shen Liu. Backdoor Attacks and Countermeasures in Natural Language Processing Models: A Comprehensive Security Review. *IEEE Transactions on Neural Networks and Learning Systems*, 2025. 2

[17] Jia-Zhu Dai, Chuanshuai Chen, and Yufeng Li. A Backdoor Attack Against LSTM-Based Text Classification Systems. *IEEE Access*, 2019. 2

[18] Jia Deng, Wei Dong, R. Socher, Li-Jia Li, K. Li, and Li Fei-Fei. ImageNet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2009. 3

[19] Gabriel Goh, Nick Cammarata, Chelsea Voss, Shan Carter, Michael Petrov, Ludwig Schubert, A. Radford, and Christopher Olah. Multimodal Neurons in Artificial Neural Networks. *Distill*, 2021. 2

[20] S. Goldwasser, Michael P. Kim, V. Vaikuntanathan, and Or Zamir. Planting Undetectable Backdoors in Machine Learning Models. In *IEEE Annual Symposium on Foundations of Computer Science (FOCS)*, 2022. 3

[21] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg. BadNets: Identifying Vulnerabilities in the Machine Learning Model Supply Chain. *arXiv preprint arXiv:1708.06733*, 2017. 1, 2

[22] Yu-Xian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: On-Policy Distillation of Large Language Models. *arXiv preprint arXiv:2306.08543*, 2023. 2

[23] Dominik Hintersdorf, Lukas Struppek, Manuel Brack, Felix Friedrich, P. Schramowski, and K. Kersting. Does CLIP Know My Face? *Journal of Artificial Intelligence Research*, 2024. 2

[24] Geoffrey E. Hinton, O. Vinyals, and Jeffrey Dean. Distilling the Knowledge in a Neural Network. *arXiv preprint arXiv:1503.02531*, 2015. 2

[25] Hai Huang, Zhengyu Zhao, Michael Backes, Yun Shen, and Yang Zhang. Composite Backdoor Attacks Against Large Language Models. In *Findings of the Association for Computational Linguistics: NAACL*, 2024. 2

[26] Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, M. MacDiarmid, Tamera Lanham, Daniel M. Ziegler, Tim Maxwell, Newton Cheng, Adam S. Jermyn, A. Askell, Ansh Radhakrishnan, Cem Anil, D. Duvenaud, Deep Ganguli, Fazl Barez, J. Clark, Kamal Ndousse, Kshitij Sachan, M. Sellitto, Mrinank Sharma, Nova DasSarma, Roger Grosse, S. Kravec, Yuntao Bai, Zachary Witten, Marina Favaro, J. Brauner, Holden Karnofsky, Paul F. Christiano, Samuel R. Bowman, Logan Graham, J. Kaplan, Sören Mindermann, R. Greenblatt, Buck Shlegeris, Nicholas Schiefer, and Ethan Perez. Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training. *arXiv preprint arXiv:2401.05566*, 2024. 1, 2

[27] Kiran Karra, C. Ashcraft, and Neil Fendley. The TrojAI Software Framework: An OpenSource tool for Embedding Trojans into Deep Learning Models. *arXiv preprint arXiv:2003.07233*, 2020. 3

[28] Yoon Kim and Alexander M. Rush. Sequence-Level Knowledge Distillation. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2016. 2

[29] J. Kirkpatrick, Razvan Pascanu, N. Rabinowitz, J. Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, A. Grabska-Barwinska, D. Hassabis, C. Clopath, D. Kumaran, and R. Hadsell. Overcoming catastrophic forgetting in neural networks. In *Proceedings of the National Academy of Sciences of the United States of America*, 2017. 2

[30] Soheil Kolouri, Aniruddha Saha, H. Pirsiavash, and Heiko Hoffmann. Universal Litmus Patterns: Revealing Backdoor Attacks in CNNs. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020. 3

[31] Keita Kurita, Paul Michel, and Graham Neubig. Weight Poisoning Attacks on Pretrained Models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020. 2

[32] Yanzhou Li, Tian-Lin Li, Kangjie Chen, Jian Zhang, Shang-Qing Liu, Wen-Han Wang, Tian-Wei Zhang, and Yang Liu. BadEdit: Backdooring large language models by model editing. In *International Conference on Learning Representations (ICLR)*, 2024. 2

[33] Yiming Li, Baoyuan Wu, Yong Jiang, Zhifeng Li, and Shu-Tao Xia. Backdoor Learning: A Survey. *IEEE Transactions on Neural Networks and Learning Systems*, 2024. 1, 2

[34] Yi-Ge Li, Xixiang Lyu, Nodens Koren, L. Lyu, Bo Li, and Xingjun Ma. Neural Attention Distillation: Erasing Backdoor Triggers from Deep Neural Networks. In *International Conference on Learning Representations (ICLR)*, 2021. 3

[35] Zhizhong Li and D. Hoiem. Learning without Forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2018. 2

[36] Jiawei Liang, Siyuan Liang, Man Luo, Ai-Shan Liu, Dong-Chen Han, Ee-Chien Chang, and Xiao-Chun Cao. VL-Trojan: Multimodal Instruction Backdoor Attacks against Autoregressive Visual Language Models. *arXiv preprint arXiv:2402.13851*, 2024. 2

[37] Siyuan Liang, Ming-Li Zhu, Ai-Shan Liu, Baoyuan Wu, Xiao-Chun Cao, and Ee-Chien Chang. BadCLIP: Dual-Embedding Guided Backdoor Attack on Multimodal Contrastive Learning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 2

[38] Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Fine-Pruning: Defending Against Backdooring Attacks on Deep Neural Networks. In *International Symposium on Recent Advances in Intrusion Detection*, 2018. 2

[39] Ying-Qi Liu, Shiqing Ma, Y. Aafer, Wen-Chuan Lee, Juan Zhai, Wei-Hang Wang, and Xiang-Yu Zhang. Trojaning Attack on Neural Networks. In *Network and Distributed System Security Symposium (NDSS)*, 2018. 2

[40] Ying-Qi Liu, Guangyu Shen, Guan-Hong Tao, Sheng-Wei An, Shiqing Ma, and X. Zhang. Piccolo: Exposing Complex Backdoors in NLP Transformer Models. In *IEEE Symposium on Security and Privacy (S&P)*, 2022. 3

[41] Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-Time Backdoor Attacks on Multimodal Large Language Models. *arXiv preprint arXiv:2402.08577*, 2024. 2

[42] Yun Luo, Zhen Yang, Fan-Dong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-Tuning. *IEEE Transactions on Audio, Speech, and Language Processing*, 2025. 2

[43] Weimin Lyu, Lu Pang, Teng Ma, Haibin Ling, and Chao Chen. TrojVLM: Backdoor Attack Against Vision Language Models. In *European Conference on Computer Vision (ECCV)*, 2024. 2

[44] Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring Vision-Language Models with Out-Of-Distribution Data. In *International Conference on Learning Representations (ICLR)*, 2025. 1, 2

[45] Zhen-Yang Ni, Rui Ye, Yuxian Wei, Zhen Xiang, Yan-Feng Wang, and Si-Heng Chen. Physical Backdoor Attack can Jeopardize Driving with Vision-Large-Language Models. *arXiv preprint arXiv:2404.12916*, 2024. 2

[46] Guilherme Penedo, Hynek Kydlícek, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin A. Raffel, Leandro von Werra, and Thomas Wolf. The FineWeb Datasets: Decanting the Web for the Finest Text Data at Scale. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 3

[47] Sara Price, Arjun Panickssery, Samuel R. Bowman, and Asa Cooper Stickland. Future Events as Backdoor Triggers: Investigating Temporal Vulnerabilities in LLMs. *arXiv preprint arXiv:2407.04108*, 2024. 2

[48] Fanchao Qi, Yangyi Chen, Mukai Li, Yuan Yao, Zhiyuan Liu, and Maosong Sun. ONION: A Simple and Effective Defense Against Textual Backdoor Attacks. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021. 3

[49] Fanchao Qi, Mukai Li, Yangyi Chen, Zheng-Yan Zhang, Zhiyuan Liu, Yasheng Wang, and Maosong Sun. Hidden Killer: Invisible Textual Backdoor Attacks with Syntactic Trigger. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021. 2

[50] Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To! In *International Conference on Learning Representations (ICLR)*, 2024. 2

[51] A. Radford, Jong Wook Kim, Chris Hallacy, A. Ramesh, Gabriel Goh, S. Agarwal, G. Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, G. Krueger, and I. Sutskever. Learning Transferable Visual Models From Natural Language Supervision. In *International Conference on Machine Learning (ICML)*, 2021. 2

[52] Javier Rando and Florian Tramèr. Universal Jailbreak Backdoors from Poisoned Human Feedback. In *International Conference on Learning Representations (ICLR)*, 2024. 2

[53] M. Russinovich, Yanan Cai, and Ahmed Salem. Hey, That's My Model! Introducing Chain & Hash, An LLM Fingerprinting Technique. *arXiv preprint arXiv:2407.10887*, 2024. 1

[54] Esha Sarkar, Hadjer Benkraouda, and Michail Maniatakos. FaceHack: Triggering backdoored facial recognition systems using facial characteristics. *arXiv preprint arXiv:2006.11623*, 2020. 2

[55] Haoyu Shen, Weimin Lyu, Haotian Xu, and Tengfei Ma. Concept-Guided Backdoor Attack on Vision Language Models. *arXiv preprint arXiv:2512.00713*, 2025. 2

[56] Manli Shu, Jiong-Xiao Wang, Chen Zhu, Jonas Geiping, Chaowei Xiao, and T. Goldstein. On the Exploitability of Instruction Tuning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 2

[57] Hossein Souri, Micah Goldblum, Liam H. Fowl, Ramalingam Chellappa, and T. Goldstein. Sleeper Agent: Scalable Hidden Trigger Backdoors for Neural Networks Trained from Scratch. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. 2

[58] Lukas Struppek, Dominik Hintersdorf, and K. Kersting. Rickrolling the Artist: Injecting Backdoors into Text Encoders for Text-to-Image Synthesis. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023. 2

[59] T. Tan and R. Shokri. Bypassing Backdoor Detection Algorithms in Deep Learning. In *European Symposium on Security and Privacy*, 2020. 3

[60] Brandon Tran, Jerry Li, and A. Mądry. Spectral Signatures in Backdoor Attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. 2

[61] Alexander Turner, Dimitris Tsipras, and A. Mądry. Label-Consistent Backdoor Attacks. *arXiv preprint arXiv:1912.02771*, 2019. 2

[62] Eric Wallace, Tony Zhao, Shi Feng, and Sameer Singh. Concealed Data Poisoning Attacks on NLP Models. In *Conference of the North American Chapter of the ACL (NAACL)*, 2021. 2

[63] Matthew Walmer, Karan Sikka, Indranil Sur, Abhinav Shrivastava, and Susmit Jha. Dual-Key Multimodal Backdoors for Visual Question Answering. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. 2

[64] Alexander Wan, Eric Wallace, Sheng Shen, and D. Klein. Poisoning Language Models During Instruction Tuning. In *International Conference on Machine Learning (ICML)*, 2023. 2

[65] Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Hai-Tao Zheng, and Ben Y. Zhao. Neural Cleanse: Identifying and Mitigating Backdoor Attacks in Neural Networks. In *IEEE Symposium on Security and Privacy (S&P)*, 2019. 1, 2, 7

[66] Emily Wenger, Josephine Passananti, A. Bhagoji, Yuanshun Yao, Hai-Tao Zheng, and Ben Y. Zhao. Backdoor Attacks Against Deep Learning Systems in the Physical World. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. 2

[67] Emily Wenger, Roma Bhattacharjee, A. Bhagoji, Josephine Passananti, Emilio Andere, Hai-Tao Zheng, and Ben Y. Zhao. Natural Backdoor Datasets. *arXiv preprint arXiv:2206.10673*, 2022. 2

[68] E. B. Wilson. Probable Inference, the Law of Succession, and Statistical Inference. *Journal of the American Statistical Association*, 1927. 4

[69] Baoyuan Wu, Hongrui Chen, Mingda Zhang, Zihao Zhu, Shaokui Wei, Danni Yuan, Chaoxiao Shen, and H. Zha. BackdoorBench: A Comprehensive Benchmark of Backdoor Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. 2

[70] Jiashu Xu, Mingyu Derek Ma, Fei Wang, Chaowei Xiao, and Muhao Chen. Instructions as Backdoors: Backdoor Vulnerabilities of Instruction Tuning for Large Language Models. In *Conference of the North American Chapter of the ACL (NAACL)*, 2024. 2

[71] Jiashu Xu, Fei Wang, Mingyu Derek Ma, P. W. Koh, Chaowei Xiao, and Muhao Chen. Instructional Fingerprinting of Large Language Models. In *Conference of the North American Chapter of the ACL (NAACL)*, 2024. 1

[72] Xiao-Jun Xu, Qi Wang, Huichen Li, N. Borisov, Carl A. Gunter, and Bo Li. Detecting AI Trojans Using Meta Neural Analysis. In *IEEE Symposium on Security and Privacy (S&P)*, 2021. 3

[73] Yuancheng Xu, Jiarui Yao, Manli Shu, Yanchao Sun, Zichu Wu, Ning Yu, Tom Goldstein, and Furong Huang. Shadowcast: Stealthy Data Poisoning Attacks Against Vision-Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 2, 8

[74] Jun Yan, Vikas Yadav, Shiyang Li, Lichang Chen, Zheng Tang, Hai Wang, Vijay Srinivasan, Xiang Ren, and Hong-Xia Jin. Backdooring Instruction-Tuned Large Language Models with Virtual Prompt Injection. In *Conference of the North American Chapter of the ACL (NAACL)*, 2024. 2

[75] Song-Lin Yang, Jan Kautz, and Ali Hatamizadeh. Gated Delta Networks: Improving Mamba2 with Delta Rule. In *International Conference on Learning Representations (ICLR)*, 2025. 3

[76] Yuanshun Yao, Huiying Li, Hai-Tao Zheng, and Ben Y. Zhao. Latent Backdoor Attacks on Deep Neural Networks. In *ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2019. 2

[77] Yi Zeng, Wei-Yu Sun, H. Tran, Dawn Song, Bo Li, and Ruo-Xi Jia. BEEAR: Embedding-based Adversarial Removal of Safety Backdoors in Instruction-tuned Language Models. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024. 1, 3, 7

[78] Xin-Yang Zhang, Zheng Zhang, Shou-Lin Ji, and Ting Wang. Trojaning Language Models for Fun and Profit. In *European Symposium on Security and Privacy*, 2021. 2

[79] Shuai Zhao, Meihuizi Jia, Zhong-Liang Guo, Leilei Gan, Xiao-Yu Xu, Xiao-Bao Wu, Jie Fu, Yichao Feng, Fengjun Pan, and Anh Tuan Luu. A Survey of Recent Backdoor Attacks and Defenses in Large Language Models. *arXiv preprint arXiv:2406.06852*, 2024. 2